# Ethical AI Strategy for Survivor-Created Records

## Draft Report
## Transparify – Team 7

**Team members**: Rohan Ramkumar, Yilin Li, Seyun Jung, Juan David Lopez, Sakada Sen, Xu Phuoc Le

**Mentor**: David Millar-Powell

**Client**: Sarah McLoughlin

# Introduction

## Project Background & Problem

### Main challenge

The main challenge is determining how to integrate AI into the products and services being offered while preserving authorship, consent, accessibility, and ethical governance of survivor-created records. The team must research AI trends, relevant technologies, and data standards to design systems that track authorship, maintain transparent metadata, and support accessible AI features such as Easy-Read summaries while protecting sensitive information.

### Why is it important?

This challenge is important because the organisation works with sensitive records and lived-experience data, which must remain accurate, traceable, and ethically managed. If AI systems are implemented without proper metadata standards, consent tracking, and governance, there is a risk of misuse, loss of authorship, ethical violations, or harm to vulnerable communities. Understanding how to implement AI responsibly will allow the company to improve accessibility and scalability while protecting trust, transparency, and the integrity of survivor-authored records.

## Project Objectives

Based on the client brief, three primary objectives have been identified and summarised to guide the team's work during the program. And to ensure that the project outcomes align with the client's expectations.

- Objective 1: Metadata / authorship standards
- Objective 2: Accessibility-focused AI features and information needs
- Objective 3: Ethics, evaluation, and staged roadmap

# Research approach and scope

To clarify the scope of this project and each member's contribution, the following Table 1 summaries the main project outcomes, the key activities to achieve them, and the activity leads responsible for each area.

| Project Outcomes | Key activities to achieve the outcome | Activity lead |
|---|---|---|
| **Recommend metadata and authorship practices for survivor-created records** | 1. Review existing metadata and authorship standards relevant to survivor-created records.<br>2. Identify the minimum information required to track | Rohan Ramkumar<br>Juan David Lopez |

| | | |
|---|---|---|
| | authorship, ownership and consent.<br>3. Recommend procedures for creating, updating, reviewing and managing records. | |
| **Evaluate and assess accessibility-focused AI opportunities and information requirements (AAC, easy read)** | 1. Research accessibility-focused AI relevant features and frameworks to help indigenous people and survivors.<br>2. List and compare the benefits, risks and limitations of each feature.<br>3. Map out the required information, safeguards and transparency measures for each proposed feature. | Yilin Li<br>Seyun Jung |
| **Develop ethical considerations and implementation recommendations for future AI use** | 1. Identify ethical, privacy, governance and legal risks involved in applying AI to sensitive survivor-related records.<br>2. Develop an ethics checklist from identified risks to assess potential AI uses and misuses.<br>3. Propose practical recommendations and a staged roadmap for future AI adoption. | Sakada Sen<br>Xuan Phuoc Le |
| **Produce a report summarising findings and practical recommendations** | 1. Consolidate research findings from the project.<br>2. Prepare a draft report for review.<br>3. Revise and finalise the report with recommendations. | All members |
| **Resources required** | ● Project brief and background materials<br>● Secondary research sources and case studies<br>● Relevant metadata / accessibility / AI governance references<br>● Client feedback or clarification where needed | |
| **Project Constraints** | ● Limited project timeframe<br>● Student-level scope and limited technical capacity<br>● Reliance on secondary research rather than primary testing or system development<br>● Limited access to client data, stakeholders, or internal | |

| | |
|---|---|
| | systems<br>● Ethical sensitivity of survivor-related records and AI use |

*Table 1. Project scope summary*

# Objective 1: Metadata / authorship standards

## Widely used authorship and metadata standards

| **Standard** | **Fields** | **Strengths** | **Limitations** | **Relevance to the project** |
|---|---|---|---|---|
| *Dublin Core Metadata Element Set* | Title, Date, Id, Publisher, Creator | Widely adopted, simple, acknowledges authorship | Limited detail on ownership and consent | Is relevant for authorship tracking |
| *Schemma.org* | Author, Date, publisher | Easy to integrate with web systems | Not designed for compliance or auditing | Useful for interoperability with web services |
| *ISO 19115* | Responsible party, Date, Id | Defines clear ownership | Domain-specific | Good for role-based ownership models |
| *Provenance Authoring Versioning* | Author, date, version | Clear distinction of roles and versions | Not widely adopted | Good for distinguishing authorship and ownership , as well as keeping track of versions |
| *Australian Government Recordkeeping Metadata Standard (*AGRkMS) | Detailed agent roles, creation date, lifecycle, consent aware access rights | Highly aligned with ethical recordkeeping and Australian law | Complex structured implementation | Good for keeping track of the lifecycle of the record. |

Based on the frameworks researched, adopting a targeted subset rather than all available standards provides a more practical and effective solution. Dublin Core Metadata Element Set can be used for capturing core metadata such as authorship, dates, and purpose, while Provenance Authoring and Versioning (PAV) enables clear tracking of edits and version history. ThenAGRkMS further strengthens this by incorporating governance, consent, and lifecycle management aligned with Australian standards. Other frameworks such as schema.org were considered but are less critical for internal recordkeeping and can be

treated as optional extensions. This approach avoids unnecessary complexity while ensuring that authorship is preserved, changes are traceable, and records are protected from misuse or loss.

## Minimum information required to track authorship, ownership and consent

To ensure compliance with Australian regulations such as the Privacy Act 1988 and the Australian Privacy Principles (APPs), as well as to protect the authorship, ownership, and consent of the records, the following data should be kept by the system.

For authorship and ownership tracking, the system should record:

- Author ID: unique identifier of the original creator of the asset
- Owner ID: unique identifier of the current owner
- Asset ID: identifier of the record or resource
- Ownership type: role associated with the asset (e.g. creator, editor, viewer)
- Timestamp of ownership creation: when ownership was established

For consent tracking, the system should store:

- User ID: identifier of the user providing consent
- Consent type: what the user is agreeing to (e.g. data collection, sharing)
- Consent status: granted or withdrawn
- Timestamp: when consent was given or revoked
- Consent method: how consent was obtained (e.g. checkbox, form)
- Policy version: the specific version of terms agreed to

Additionally, the system should maintain an audit trail of changes to ownership and consent to ensure traceability and accountability over time, aligning with best practices in secure information management.

## Proposed Workflows: Record Creation and Version Control

### Workflow 1: Independent Record Creation & Export (Online)

*Designed for self-advocates securely logging their own experiences with reliable internet access.*

- **Create:** The author opens a simple, accessible digital form and fills in their story. In the background, the system automatically converts this into a structured document and saves it to a secure, version-controlled GitHub branch. The system automatically logs their Author ID, the exact timestamp, the record's purpose, and their active Consent Status.

- **Update:** As the author adds details over time, they simply click "Save." The system automatically registers each save as a new GitHub commit, creating an immutable, mathematical ledger of exactly what was changed, at what time, and by whom, preventing invisible edits or erasure.

- **Review:** When the record is ready, the author clicks a "Submit for Review" button. The system automatically generates a GitHub Pull Request and notifies a trusted community reviewer. The reviewer checks the system dashboard specifically to ensure the minimum metadata checklist is complete and the consent status is clearly marked.

- **Export:** Upon the reviewer's approval, the record is finalized. An automated script is triggered to export the document into a clean, read-only PDF or Easy-Read format. The exported document automatically appends a "Metadata Certificate" at the bottom, verifying the file's authorship and complete chronological history.

## Workflow 2: Assisted Creation Capturing Dual-Authorship

*Designed for survivors utilizing AAC or Easy-Read methods alongside a community advocate, ensuring the survivor's primary authorship is not erased by the assistant.*

- **Create:** The advocate and survivor co-create the record using an accessible interface. The interface prompts them to assign roles. In the background, the system explicitly logs the survivor as the "Primary Author" and the advocate as the "Assisting Author," and records that consent was gathered via AAC methods.

- **Update:** If the survivor requests amendments, the advocate updates the file and checks a required box confirming the survivor's approval. The system's background commit message automatically explicitly states the nature of the change and logs the verified AAC approval for that specific edit.

- **Review:** A secondary reviewer examines the submitted record. Their sole focus is verifying that the dual-authorship metadata aligns with the Alliance's ethics checklist, ensuring the advocate has not overstepped or failed to log the survivor's explicit consent for the updates.

- **Export:** Once approved, the automated export process generates an accessible format (e.g., a plain-language summary). The system specifically formats the header to prominently display both the Primary Author and the Assisting Author, ensuring systemic memory of the collaborative effort.

## Workflow 3: Offline Update & Sync (Low Connectivity)

*Designed for remote locations or situations with unstable internet, ensuring data integrity is maintained when offline.*

- **Create / Update (Offline):** An advocate creates or edits a record using an offline-capable Alliance application on their device. As they save their work, the app's internal system locally tracks all changes, capturing the exact offline timestamps and the local author’s credentials without needing an internet connection.

- **Review (Syncing Phase):** Once the device connects to the internet, the advocate clicks "Sync." The app securely pushes the local data to the central repository and submits the record for review. The reviewer can now see the chronological history of all edits made during the offline period, exactly as they occurred.

- **Export:** The reviewer verifies the integrity of the offline timestamps and metadata. Once finalized, the system tags the record with a "Synced" status and securely exports it to the relevant allied health or legal advocate database, maintaining the verifiable trail back to the original offline creation point.

# Objective 2: Accessibility-focused AI features and information needs

## Purpose and approach

This goal aims to study practical artificial intelligence (AI) features with accessibility as the primary target users. These features need to support the client in handling records written by survivors while ensuring security, traceability, and transparency. And our client has explicitly requested that we review existing relevant AI tools and use reliable public resources to determine their functional requirements. Finally, we will map each function to the information required for its safe operation, and then compare two practical solutions which are suitable for unstable network access or poor network conditions.

This section focuses on the research and evaluation of accessibility-focused AI features across several key categories, including plain-text simplification and comprehension support, symbol-supported communication, audio and video accessibility, image accessibility, and accessible extraction from scanned records. The reason for choosing these features is that they closely align with the core needs explicitly stated by the client. Therefore, we will conduct searches and screenings around AI tools that propose these features.

## Review of accessibility-focused AI examples and feature requirements

The sources indicate that accessibility-oriented AI is most valuable in enhancing information understanding, communication, and access, provided that it does not replace the original records or obscure their creators. The functionalities evaluated in this project cover several broader areas: text accessibility and simplification, augmentative and alternative communication (AAC) systems and symbol-supported communication, audio and video accessibility, image accessibility, and the accessibility of scanned or non-editable records. Public guidelines demonstrate that each area has its unique functional requirements. For example, text-based features should be audience-centered, clear, and easy to understand; AAC tools should be personalised and flexible; audio and video features should prioritise accuracy and completeness; image-related features should reflect the purpose and context of the visual content; and tools based on optical character recognition (OCR) should ensure accurate and accessible extraction of text from scanned materials.

# Feature review, information, and transparency table

| Feature ID | Proposed AI feature | Target accessibility group | What the feature does | Functional requirements supported by public sources | Potential benefit | Main risks/limitations | Suggested proiority | Sources |
|---|---|---|---|---|---|---|---|---|
| F1 | Plain-language summary generator | People under pressure, with low literacy, neurodivergent users, and advocates | Produces a shorter version of an article in everyday language | Identify the target audience; highlight the main information; use short paragraphs and common vocabulary, and avoid using jargon. | Improve the comprehension of the initial reading without altering the original document. | Summaries may overlook subtle differences, leading to user misunderstandings. | High | Digital.gov — Plain Language: Principles CDC — Plain Language |
| F2 | Easy-Read generator | People with intellectual disability, cognitive disability, low literacy, limited English proficiency, or users reading under stress | Rewrite source into "Easy-Read" using simpler words, shorter sentences, clearer layout, etc. | Emphasize using simple vocabulary, short sentences, expressing one idea at a time, using positive wording, maintaining consistent formatting, and having members of the target audience proofread the text. | Best fit with AAC/Easy-Read and existing prototype; high strategic alignment | Similar to F1, if auto-generated badly, Easy-Read can become inacurrate or even risks oversimplifying | High | Inclusion Europe — Information for All: European Standards for Making Information Easy to Read and Understand Inclusion Europe — Easy-to-Read |

| | | | | | | | | |
|---|---|---|---|---|---|---|---|---|
| F3 | AAC / symbol-supported communication-card generator | People communicate with symbols, boards or speech-generating AAC systems | Extract key message and then translate into symbol cards, prompts or phrase sets | AAC emphasises that vocabulary should be functional, personalised, and relevant to an individual's language, dialect, culture, and identity. | The client has already worked with AAC cards | High personalisation risks may face to. The choice of symbols, vocabulary, and cultural fit cannot guarantee the degree of generalisation. Therefore, human-machine collaboration (with some symbol selections requiring human verification) is very necessary here. | Medium | ASHA Practice Portal — Augmentative and Alternative Communication |
| F4 | Auto-captioning and transcript generation | Deaf and hard-of-hearing users | Create captions and transcripts for recorded interviews, webinars, videos, etc. | WCAG requires that pre-recorded synchronised media must be accompanied by captions; The FCC's quality standards emphasise accuracy, timeliness, completeness, and placement. | Useful to support video-based resources | Speech recognition errors may lead to misstatements in reports; some sensitive words might be recognised directly without undergoing review. | Medium | W3C — Web Content Accessibility Guidelines (WCAG) 2.2 FCC — Closed Captioning on Television |
| F5 | Alternative-text generator | Blind or low-vision | Generates text alternatives for | The W3C guidelines | Work as an important | Generative systems may | Medium | W3C WAI — Images |

| | | | | | | | | |
|---|---|---|---|---|---|---|---|---|
| | | people, and anyone who cannot load or see images | images, screenshots, diagrams, and figures. | state that alternative text should reflect the purpose of the image, rather than just listing visible objects; unless the image itself is very complex and requires a more detailed description elsewhere, concise text should be prioritised. | supporting feature | produce false content, ignore context, or leak sensitive details. | | Tutorial W3C WAI — An alt Decision Tree |
| F6 | Audio-description generator | Blind or low-vision users, and users who cannot fully interpret visual content in videos | Creates spoken descriptions of visual information in videos, such as actions, scene changes, and on-screen text. Could be appliable to both pre-recorded or live content. | Audio description should cover important visual content that is not already available in the main soundtrack, and it should be synchronized with the video according to W3C. | Work as an important supporting feature for video-based resources | Auto-description can miss emotional context, misinterpret, or expose sensitive visual details without enough human review. | Medium | W3C WAI — Understanding Success Criterion 1.2.5: Audio Description (Prerecorded) |
| F7 | OCR and accessible text extraction for scanned records | Screen-reader users, blind/low-vision users, and anyone working with scanned PDFs | Converts scanned documents into machine-readable text so they can be | W3C notes that scanned image-only PDFs are inherently inaccessible because | Is necessary to enable many features proposed in this document. | OCR can misread names, dates, and sensitive details, so outputs should ideally be | High | W3C — PDF7: Performing OCR on a Scanned PDF Document to Provide Actual Text |

| | | | | | | | | |
|---|---|---|---|---|---|---|---|---|
| | | or image-based records. | searched, selected, read aloud, or reformatted into summaries/Easy-Read. | assistive technologies cannot read or extract the words, and recommends OCR when the original text source is unavailable. | | checked before use. | | |
| F8 | Caption cleanup / speaker-label enhancer | Deaf and hard-of-hearing users. | Improves raw auto-captions by adding speaker labels, punctuation, and meaningful non-speech sounds. | W3C's captions guidance says captions should not only include dialogue but also identify who is speaking and include meaningful sound information. | Improvement to F4 to be more specific and practical for recorded interviews, webinars, and survivor-story media with multiple speakers. | Mislabelled speakers or missing sound cues could change meaning, especially in sensitive testimony. | Medium | W3C WAI — Understanding Success Criterion 1.2.2: Captions (Prerecorded) |
| F9 | Plain-language glossary | People under pressure, users with low literacy, neurodivergent users, and users unfamiliar with legal or technical language. | Detects difficult words, acronyms, and specialist terms and gives short plain-language explanations inline or on demand. | Plain-language guidance emphasizes writing for the audience, using clear language, and avoiding jargon; Easy-Read guidance also says difficult words should be explained clearly. | Is an efficient substitute to completely rewriting documents. | Definitions may oversimplify or misinterpret technical terms. | Medium | Digital.gov — Plain Language: Principles |

## Discussion of feature priorities

Based on the feature review table, the highest-priority features are plain-language summary generation, Easy-Read generation, and OCR with accessible text extraction. These should be prioritised first because they offer the widest immediate benefit across the client's likely use cases while also supporting the broader accessibility and metadata goals of the project. Plain-language summaries and Easy-Read outputs directly align with the organisation's focus on accessible formats, including AAC and Easy-Read, and are especially relevant for users who may be under stress, have low literacy, or experience cognitive barriers when reading complex records. OCR is also a high-priority enabling feature, because scanned or image-based records must first be converted into machine-readable text before they can be summarised, rewritten into Easy-Read, or further processed in a transparent way. In practice, this makes OCR a foundational capability rather than just an optional extra.

A second tier of medium-priority features includes AAC or symbol-supported communication-card generation, auto-captioning and transcript generation, alternative-text generation, audio-description generation, and caption enhancement. These features are still valuable, but their benefits are either more specialised or more dependent on particular content types such as images, videos, or media resources. AAC card generation is strategically important because it strongly matches the client's existing AAC focus and product direction, but it also carries higher personalisation and cultural-fit risks, meaning it requires more careful human review than simpler text-accessibility features. Captioning, audio description, and alternative text are all useful extensions for multimodal accessibility, particularly if the organisation continues producing videos, interviews, webinars, toolkits, or public-facing explainers.

Overall, the feature priorities suggest a practical implementation order. The client would benefit most from first adopting features that improve access to written material and reduce comprehension barriers across many records, especially where those features can be combined with clear metadata, consent controls, and source tracing. More specialised multimodal features should follow once the organisation has established a strong foundation for accessible text processing, transparent outputs, and human review processes.

## Comparison of technical approaches

Two practical technical approaches are especially relevant to this project: an offline-first architecture and a low-connectivity cloud-assisted architecture. The client brief explicitly asks for a comparison between these two approaches for situations where internet access may be unstable or limited. Because the organisation works with sensitive survivor-authored records and aims to preserve authorship, consent, and auditability, the technical choice should not be based only on convenience or model performance. It should also reflect privacy, transparency, sustainability, and the need to avoid unnecessary vendor lock-in.

| Approach | Main strengths | Main limitations | Best-fit use cases |
|---|---|---|---|
| Offline-first | Better privacy by default; more resilient in poor or unstable internet conditions; stronger local control over records and consent-sensitive materials | Harder to maintain and update; limited by device performance; may require more local technical support | Sensitive survivor-authored records, field/community use, low-bandwidth environments, basic OCR and text-accessibility workflows |

| | | | |
|---|---|---|---|
| Low-connectivity cloud-assisted | More scalable; easier to update and improve centrally; can support stronger AI models and richer multimodal features | Depends more on internet access; raises greater consent, governance, and vendor-dependence risks; needs stricter transparency controls | Mixed environments where some connectivity is available, staged pilots, advanced summarisation or multimodal accessibility features |

On balance, the comparison suggests that neither approach is sufficient on its own for all situations. Offline-first is safer and more aligned with the client's ethics when handling highly sensitive records, while low-connectivity systems are more practical for scaling and improving feature quality over time. The strongest option is therefore a staged or hybrid model in which the most sensitive and essential functions remain available locally, while higher-powered optional features can be accessed through carefully governed remote services when appropriate. This better matches the client's emphasis on accessibility, governance, and long-term sustainability.

## Recommendations

Based on the feature review and technical comparison, three recommendations can be made for Objective 2. First, the client should prioritise text-based accessibility features with broad immediate value, especially plain-language summaries, Easy-Read generation, and OCR-enabled text extraction. These features align most closely with the organisation's existing focus on accessible records and are likely to produce the greatest practical benefit for users across multiple contexts. They also create the foundation needed for later accessibility functions by ensuring records can first be read, interpreted, and transformed in transparent ways.

Second, the client should adopt a human-review model for all accessibility outputs that affect meaning. Because the records involved are sensitive and survivor-authored, AI outputs should support access and comprehension without silently changing authorship, intent, or context. This is especially important for Easy-Read conversion, summarisation, AAC symbol selection, captions, and image descriptions, where inaccuracies may distort meaning or introduce harm. In practice, AI should be treated as a drafting and support tool rather than an autonomous decision-maker. This recommendation is strongly consistent with the brief's warning against black-box systems and uncontrolled use of sensitive participant data.

Third, the client should favour a hybrid technical approach led by offline-first principles. Core record access, metadata visibility, consent information, and basic accessibility workflows should remain available locally wherever possible, while more advanced AI processing can be layered on through low-connectivity services only when clearly justified and appropriately governed. This would allow the organisation to balance privacy, resilience, accessibility, and future scalability without committing too early to one external platform. It would also reduce the risk of vendor lock-in while still leaving room for experimentation with summarisation, Easy-Read, and multimodal tools as the organisation's capacity grows.

# Objective 3: Ethics, evaluation, and staged roadmap

## Introduction

AI tools like chatbots and Augmentative and Alternative Communication (AAC) systems can make services more accessible and help survivors find early-stage information. But bringing AI into the space of sensitive, survivor-created records comes with heavy ethical, privacy, governance, and legal risks. The reality is that AI isn't naturally equipped to provide trauma-informed care. Chatbots can't read a person's body language, assess their emotional state, or understand the full context of their situation (Mielismäki & Husso, 2025). Relying on them risks trivialising complex, potentially life-threatening domestic violence scenarios. If the organisation wants to integrate these tools safely, we have to look closely at these intersecting risks so that vulnerable communities aren't put in further danger.

## Ethical Risks

Bringing AI into survivor support introduces deep ethical challenges that directly impact user autonomy and safety. For a start, chatbots simply can't handle comprehensive situational assessments or build a real support network. This strips away the human context, reducing complex trauma to automated outputs and data points (Mielismäki & Husso, 2025; OVIC; Omoyemi, 2024). There is also a very real danger that the very tools meant to help survivors can be weaponised by perpetrators, who might hack devices or monitor online activity to exert further control (Mielismäki & Husso, 2025). On the development side, AI training data often carries the cultural and societal biases of its creators, which can end up reinforcing stereotypes or pushing marginalised victim-survivors even further to the fringes (Mielismäki & Husso, 2025; Omoyemi, 2024; OVIC). AI guidance can also become manipulative, subtly threatening a user's ability to make free, independent choices (Mielismäki & Husso, 2025). We also have to consider the burden these systems place on the user: when a survivor interacts with a chatbot, the heavy lifting of assessing the danger and maintaining the connection falls entirely on them (Mielismäki & Husso, 2025). Finally, treating deeply rooted structural problems as issues that can be "fixed" with an app or a chatbot is a form of technological solutionism. It risks distracting us from the actual work of strengthening human-led support systems (Mielismäki & Husso, 2025).

## Privacy Risks

AI systems require massive amounts of sensitive personal data to function, creating severe privacy vulnerabilities. When victim-survivors share intimate details with a chatbot, they rarely know exactly how that information is being stored, shared, or protected (Mielismäki & Husso, 2025; Omoyemi, 2024). Cloud-based AAC systems carry similar risks, exposing user data to potential breaches (Omoyemi, 2024). This vulnerability is easily exploited by perpetrators; if an abuser installs spyware or gains access to a device, help-seeking tools quickly turn into surveillance instruments (Mielismäki & Husso, 2025). Anonymity is another double-edged sword. While being anonymous might make a survivor feel safer disclosing abuse initially, it leaves no formal case record. This means repeated incidents of violence stay completely off the radar of protection systems

(Mielismäki & Husso, 2025). By trying to avoid documentation, survivors might unknowingly put themselves and their children in greater danger because critical safety information is never flagged (Mielismäki & Husso, 2025). Automated systems can also trick users into over-disclosing. Because a chatbot feels conversational and non-judgmental, survivors might share much more than they intended to (Mielismäki & Husso, 2025; OVIC). Behind the scenes, AI is incredibly good at inferring hidden details, like someone's mental health status or immigration background, from seemingly unrelated data points (OVIC). This blurs the line of what is actually considered "personal information," and raises the risk that data collected for a simple support chat gets repurposed without the survivor's knowledge or consent (OVIC).

## Governance Risks

AI is evolving so quickly that it's leaving existing governance and regulatory frameworks behind (OVIC). Right now, if an AI is the primary point of contact, there is no designated human responsible for taking concrete action, organising support, or building a safety network (Mielismäki & Husso, 2025). If a chatbot gives bad advice in a life-threatening situation or fails to spot a crisis, there is no clear accountability (Mielismäki & Husso, 2025). This creates a massive, dangerous gap for things like child protection, which might just slip through the cracks of an automated tool (Mielismäki & Husso, 2025). It doesn't help that chatbot quality varies wildly from one provider to the next, and there are no standard rules for judging their reliability in domestic violence situations (Mielismäki & Husso, 2025). We also can't assume that practitioners are ready to use these tools properly; in Finland, for instance, over half of police and social workers don't currently use any tools at all to map domestic violence (Mielismäki & Husso, 2025). If we aren't careful, deploying AI will just mirror, or even worsen, the physical service system's existing flaws in a digital space (Mielismäki & Husso, 2025). Ultimately, the survivors who actually rely on these tools are almost never involved in designing them (OVIC; Omoyemi, 2024). When you combine that exclusion with the high costs and digital literacy needed to use the tech, it creates a harsh digital divide that stops many people from getting the help they need (Omoyemi, 2024).

## Legal Risks

Blending AI automation with strict legal obligations creates a minefield of compliance issues for the organisation. For example, a survivor chatting with a bot to seek help is not the same thing as officially reporting violence. This disconnect creates serious legal gray areas around mandatory reporting and statutory child protection duties (Mielismäki & Husso, 2025). Allowing users to remain completely anonymous pits the survivor's right to autonomy against the authorities' legal obligation to document disclosures and take action (Mielismäki & Husso, 2025). If an emergency isn't documented because the user only spoke to a bot, the organisation could be failing its legal duty of care (Mielismäki & Husso, 2025). AI is also prone to misinterpreting slang, non-standard language, or cultural nuances. If a chatbot misunderstands a user and misses a critical risk assessment or makes a bad referral, the organisation opens itself up to severe liability (Mielismäki & Husso, 2025). Trust and consent are also legally binding concepts here. Users need to know immediately if they are talking to a machine. Failing to make that clear can violate informed consent laws (Mielismäki & Husso, 2025; OVIC). And legally speaking, a user clicking "agree" during a moment of crisis probably doesn't count as valid consent for their data to be used in secondary AI training later down the line (OVIC). Finally, if these AI tools aren't built with strong

multilingual support, non-English speakers will face unequal access (Mielismäki & Husso, 2025; Omoyemi, 2024). If an automated system produces discriminatory outcomes, the organisation could face serious exposure under both privacy and anti-discrimination laws (OVIC; Omoyemi, 2024).

## Ethic checklist

| **Category** | **Key questions** | **Pass criteria** | **Fail Indicator** |
|---|---|---|---|
| Consent & User Control | Has the user agreed to this specific AI use? | Explicit, informed, and revocable consent is recorded<br><br>Parental/guardian consent obtained and documented for minors or vulnerable individuals (Fedele et al., 2024)<br><br>Consent is not reused beyond its original stated purpose without re-authorisation (Fedele et al., 2024) | Consent is unclear, missing, or reused beyond original purpose<br><br>No process exists to verify guardian consent for automated decisions involving minors |
| Authorship & Voice Integrity | Does the AI preserve the survivor's voice and meaning? | Original content is linked, no distortion, edit history tracked<br><br>Author remains identifiable and traceable | Meaning is altered, simplified incorrectly, or author is not traceable |
| Privacy & Data Protection | Is sensitive data secure and protected? | Data is anonymised or access-controlled, no unsafe third-party sharing<br><br>DPIA is completed before processing any data (Fedele et al., 2024)<br><br>Privacy-by-design and data minimisation are | Identifiable data shared with external systems without safeguards<br><br>No data protection impact assessment before processing (Fedele et al., 2024)<br><br>Privacy controls added after system |

| | | | |
|---|---|---|---|
| | | applied from the start (Fedele et al., 2024)<br><br>Sensitive attributes are identified and given extra protection (Fedele et al., 2024)<br><br>Pseudonymisation, encryption, and differential privacy prevent re-identification (Fedele et al., 2024)<br><br>Data handling and security compliance are verified (QAO, 2025) | design, not built-in (Fedele et al., 2024)<br><br>Vendor data handling obligations are undefined or unverified (QAO, 2025) |
| Human Oversight | Is a humaning reviewing AI outputs? | Humans review all sensitive AI outputs<br><br>Automated decisions allow review, explanation, and challenge (Fedele et al., 2024)<br><br>A reporting mechanism exists for affected individuals, even if unaware their data was processed (QAO, 2025) | Outputs are fully automated with no human oversight<br><br>Affected individuals cannot request review or explanation of decisions (Fedele et al., 2024)<br><br>No complaints or adverse-event pathway exists (QAO, 2025) |
| Safety & Harm Prevention | Could this output cause harm or mislead? | Outputs checked for trauma sensitivity and accuracy<br><br>Vulnerable individuals treated as a distinct risk group in ethical assessments (QAO, 2025)<br><br>Staff trained on what sensitive data must be entered into AI systems (QAO, 2025) | Outputs could be harmful, misleading, or insensitive<br><br>Vulnerable-group impacts not assessed separately (QAO, 2025)<br><br>Staff lack guidance on handling sensitive personal info (QAO, 2025) |

| | | | |
|---|---|---|---|
| Bias & Fairness | Are outputs free from discriminatory patterns? | Training data bias is continuously monitored, tracking sensitive features (gender, ethnicity, health)<br>(Fedele et al., 2024)<br><br>Explainable AI is used to detect hidden discriminatory patterns<br>(Fedele et al., 2024)<br><br>Monitoring controls are active and reviewed at set intervals to catch misuse<br>(QAO, 2025) | No ongoing bias monitoring for sensitive demographics<br>(Fedele et al., 2024)<br><br>Black-box outputs not checked for discrimination<br>(Fedele et al., 2024)<br><br>Monitoring intervals are undefined or ad hoc<br>(QAO, 2025) |
| Transparency & Disclosure | Do users know they are interacting with AI? | AI clearly labelled, source and timestamps visible | Users cannot distinguish AI from human or verify outputs |
| Accessibility (AAC & Easy Read) | Is the output usable for all users? | Meets Easy Read standards and works with AAC tools | Outputs are complex, unclear, or inaccessible to target users |
| Accountability & Auditability | Can actions and outputs be traced? | Full logging of inputs, outputs, edits, and timestamps<br><br>Clear accountability assigned to a responsible person or body<br>(QAO, 2025)<br><br>Independent audits verify vendor and system compliance<br>(QAO, 2025) | No audit trail or unclear responsibility<br><br>No named accountable person for AI outcomes<br>(QAO, 2025)<br><br>Vendor compliance assumed, not verified<br>(QAO, 2025) |

## Practical recommendations

To ensure the ethical and safe use of AI in handling sensitive survivor-related records,

Roadmap

| Stage | Timeline | Goals | Key Actions | Priority Focus |
|---|---|---|---|---|
| Short Term | 0-3 months | Establish foundations and control risk | - Implement AI ethics checklist<br>- Define metadata standards (authorship, consent, timestamps)<br>- Pilot low-risk AI tools (Easy Read summarisation)<br>- Human-in-the-loop review for all outputs<br>- Secure data environments | Safety, ethics, transparency, consent |
| Medium Term | 3-9 months | Scale responsibility and strengthen governance | - Automate metadata tracking and audit logs<br>- Introduce evaluation metrics (fairness, accessibility, transparency)<br>- Expand AI outputs for AAC and Easy Read formats<br>- Test multiple AI providers to avoid vendor lock-in<br>- Strengthen governance processes | Accessibility, accountability, risk mitigation |
| Long Term | 9-18 months | Develop sustainable and transparent AI ecosystem | - Implement real-time transparency (audit trails, explainability)<br>- Establish multi-vendor/modular AI architecture<br>- Continuous monitoring for bias, safety, and accessibility<br>- Integrate survivor and community input into system design<br>- Ensure ongoing improvement and legal compliance | Sustainability, trust, long-term accountability |

# References

[try to use same format, such as APA or IEEE]


American Speech-Language-Hearing Association. (n.d.). Augmentative and alternative communication. https://www.asha.org/practice-portal/professional-issues/augmentative-and-alternative-communication/

Centers for Disease Control and Prevention. (n.d.). Plain language. https://www.cdc.gov/health-literacy/php/develop-materials/plain-language.html

Digital.gov. (n.d.). Plain language: Principles. https://digital.gov/guides/plain-language/principles

Dublin Core Metadata Initiative. (n.d.). *Dublin Core metadata element set, version 1.1*. https://www.dublincore.org/specifications/dublin-core/dces/

Fedele, A., Punzi, C., & Tramacere, S. (2024). The ALTAI checklist as a tool to assess ethical and legal implications for a trustworthy AI development in education. *Computer Law & Security Review*, 53, 105986. https://doi.org/10.1016/j.clsr.2024.105986

Federal Communications Commission. (n.d.). Closed captioning on television. https://www.fcc.gov/consumers/guides/closed-captioning-television

International Organization for Standardization. (2014). *ISO 19115-1:2014 geographic information — Metadata — Part 1: Fundamentals*. https://www.iso.org/standard/53798.html

Inclusion Europe. (2017). Information for all: European standards for making information easy to read and understand. https://www.inclusion-europe.eu/wp-content/uploads/2017/06/EN_Information_for_all.pdf

Inclusion Europe. (n.d.). Easy-to-read. https://www.inclusion-europe.eu/easy-to-read/

Mielismäki, H., & Husso, M. (2025). Ethical Implications of AI-Driven Chatbots in Domestic Violence Support. *Social Inclusion*, *13*(S6). https://doi.org/10.17645/si.9998

National Archives of Australia. (2015). *Australian Government recordkeeping metadata standard (AGRkMS)*.https://www.naa.gov.au/

Queensland Audit Office [QAO]. (2025). *Managing the ethical risks of artificial intelligence, Report 2: 2025–26.* https://www.qao.qld.gov.au/sites/default/files/2025-09/Managing%20the%20ethical%20risks%20of%20artificial%20intelligence%20%28Report%202%20%E2%80%93%202025%E2%80%9326%29.pdf

Office of the Victorian Information Commissioner [OVIC]. (n.d.). *Artificial Intelligence and Privacy – Issues and Challenges.*

https://ovic.vic.gov.au/privacy/resources-for-organisations/artificial-intelligence-and-privacy-issues-and-challenges/

Omotayo, O., (2024). Ethical Implications of AI-Driven AAC Systems: Ensuring Inclusivity and Equity in Assistive Technologies. *World Journal of Advanced Research and Reviews*, 24(1), pp. 2576-2581. https://doi.org/10.30574/wjarr.2024.24.1.3310

Schema.org. (n.d.). *Schema.org vocabulary*. https://schema.org/

World Wide Web Consortium. (n.d.). Web Content Accessibility Guidelines (WCAG) 2.2. https://www.w3.org/TR/WCAG22/

World Wide Web Consortium Web Accessibility Initiative. (n.d.). Images tutorial. https://www.w3.org/WAI/tutorials/images/

World Wide Web Consortium Web Accessibility Initiative. (n.d.). An alt decision tree. https://www.w3.org/WAI/tutorials/images/decision-tree/

World Wide Web Consortium Web Accessibility Initiative. (n.d.). Understanding Success Criterion 1.2.5: Audio description (prerecorded). https://www.w3.org/WAI/WCAG21/Understanding/audio-description-prerecorded.html

World Wide Web Consortium. (n.d.). PDF7: Performing OCR on a scanned PDF document to provide actual text. https://www.w3.org/TR/WCAG20-TECHS/PDF7.html

World Wide Web Consortium Web Accessibility Initiative. (n.d.). Understanding Success Criterion 1.2.2: Captions (prerecorded). https://www.w3.org/WAI/WCAG22/Understanding/captions-prerecorded.html

# Appendix

[put all your tables & sources with link in here]